
Taking derivative by convolution

Partial derivatives with convolution

For 2D function $f(x,y)$, the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

For discrete data, we can approximate using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

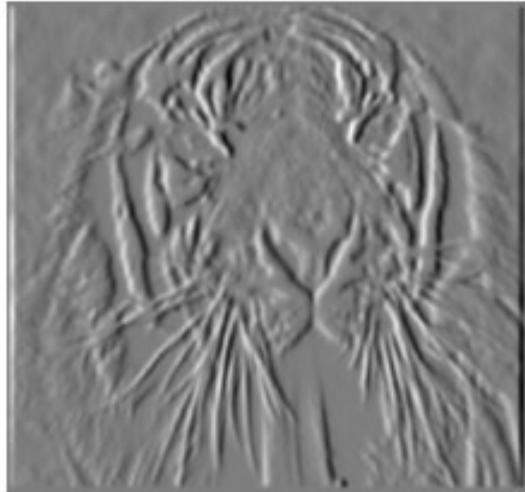
To implement above as convolution, what would be the associated filter?

Partial derivatives of an image

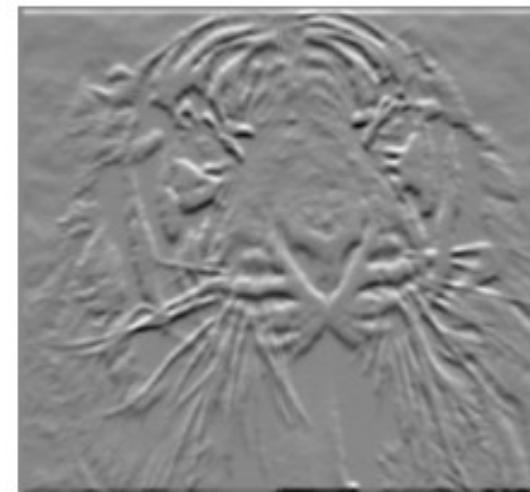


$$\frac{\partial f(x, y)}{\partial x}$$

-1	1
----	---



$$\frac{\partial f(x, y)}{\partial y}$$



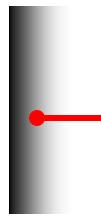
-1	1
1	-1

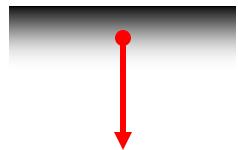
or

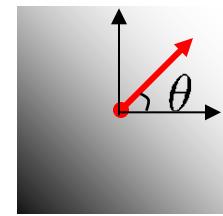
Which shows changes with respect to x?

Image gradient

The gradient of an image: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$


$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$


$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$


$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The gradient points in the direction of most rapid increase in intensity

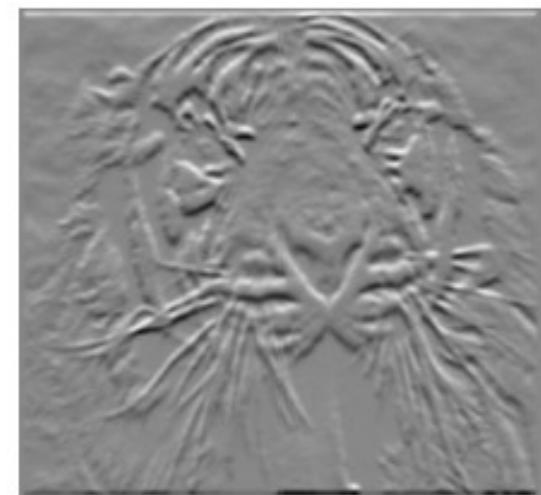
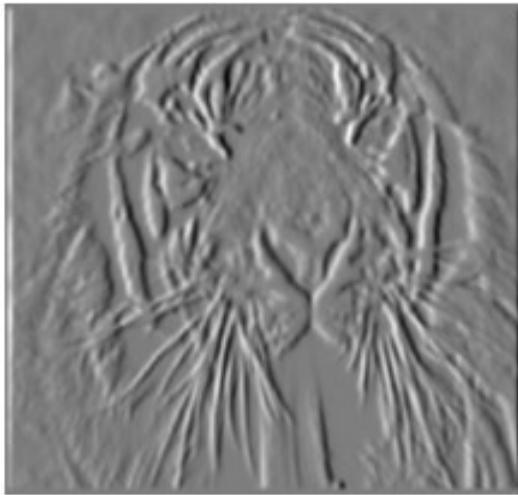
- How does this direction relate to the direction of the edge?

The gradient direction is given by $\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$

The *edge strength* is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x} \right)^2 + \left(\frac{\partial f}{\partial y} \right)^2}$$

Image Gradient



$$\frac{\partial f(x, y)}{\partial x}$$

$$\frac{\partial f(x, y)}{\partial y}$$

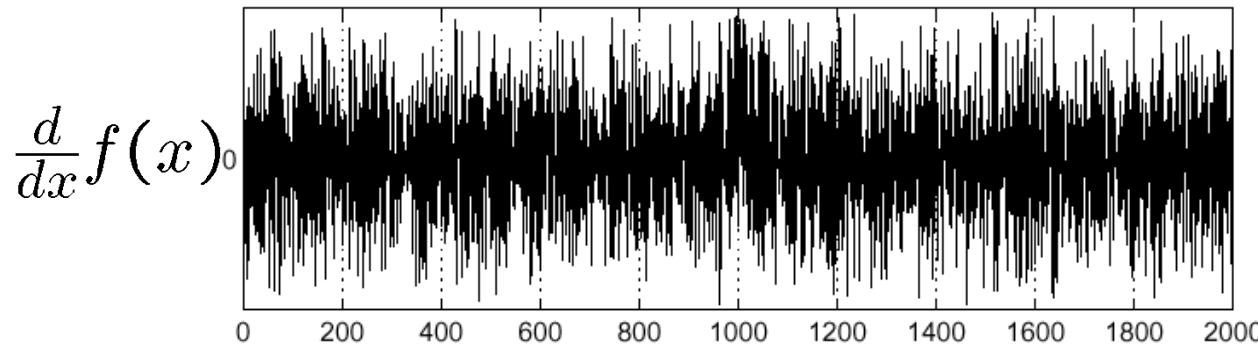
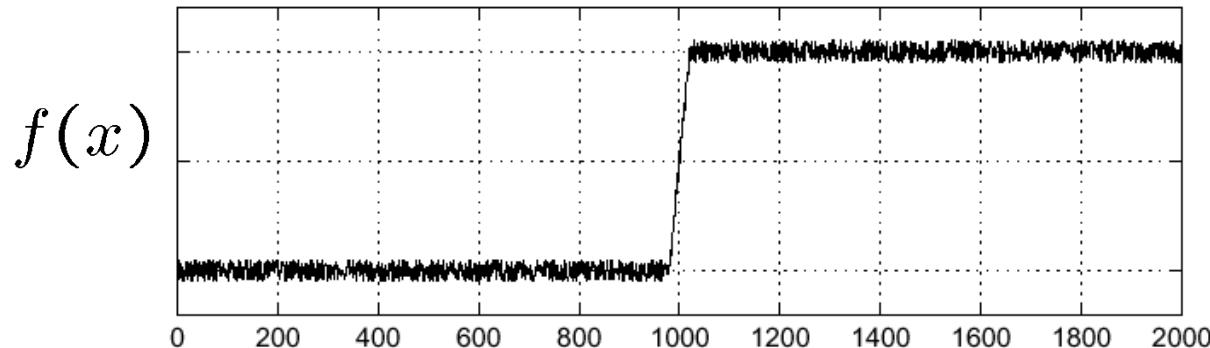


$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

Effects of noise

Consider a single row or column of the image

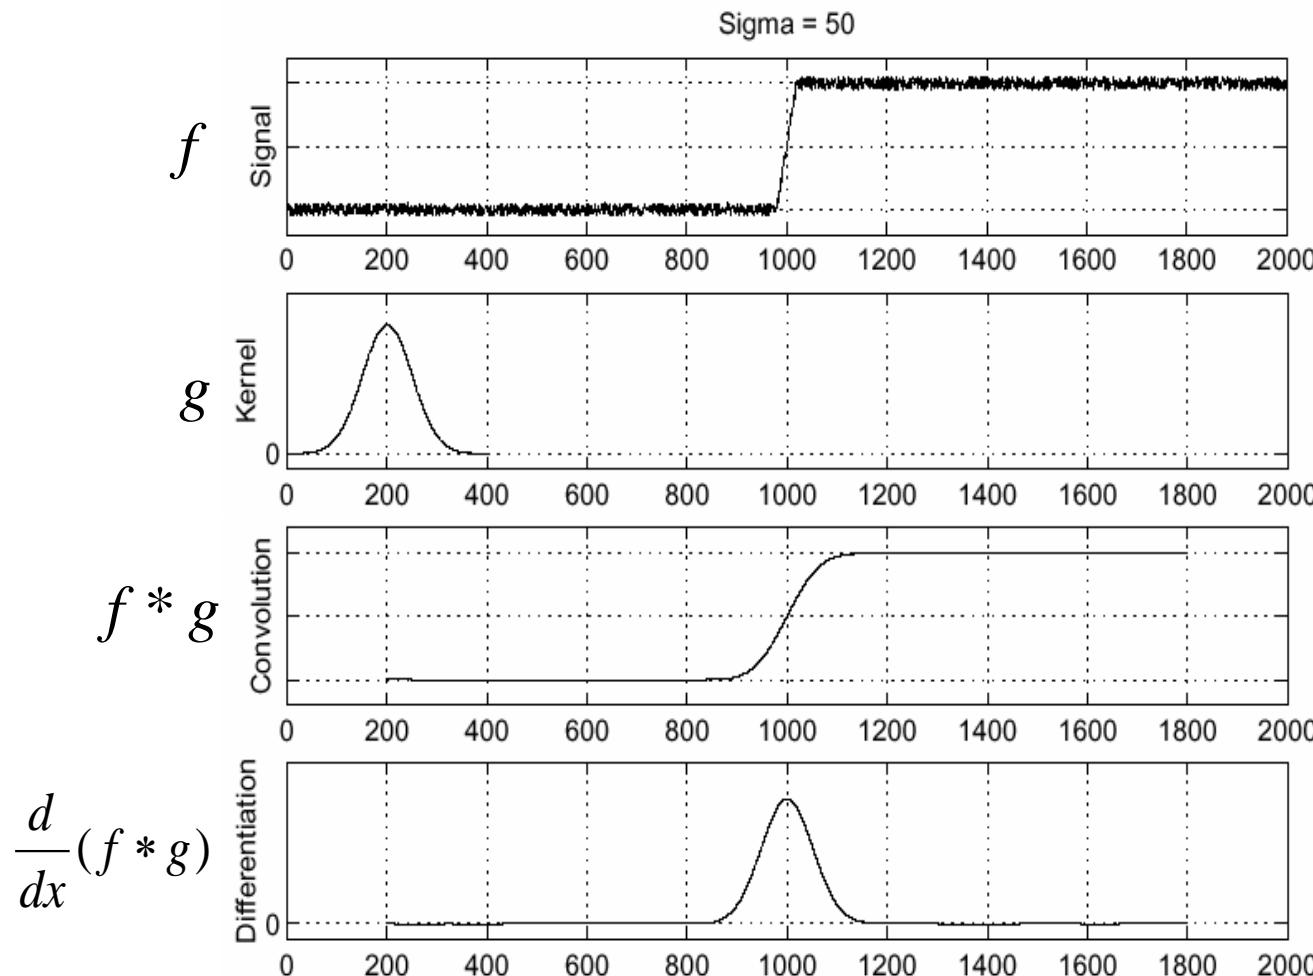
- Plotting intensity as a function of position gives a signal



Where is the edge?

Source: S. Seitz

Solution: smooth first



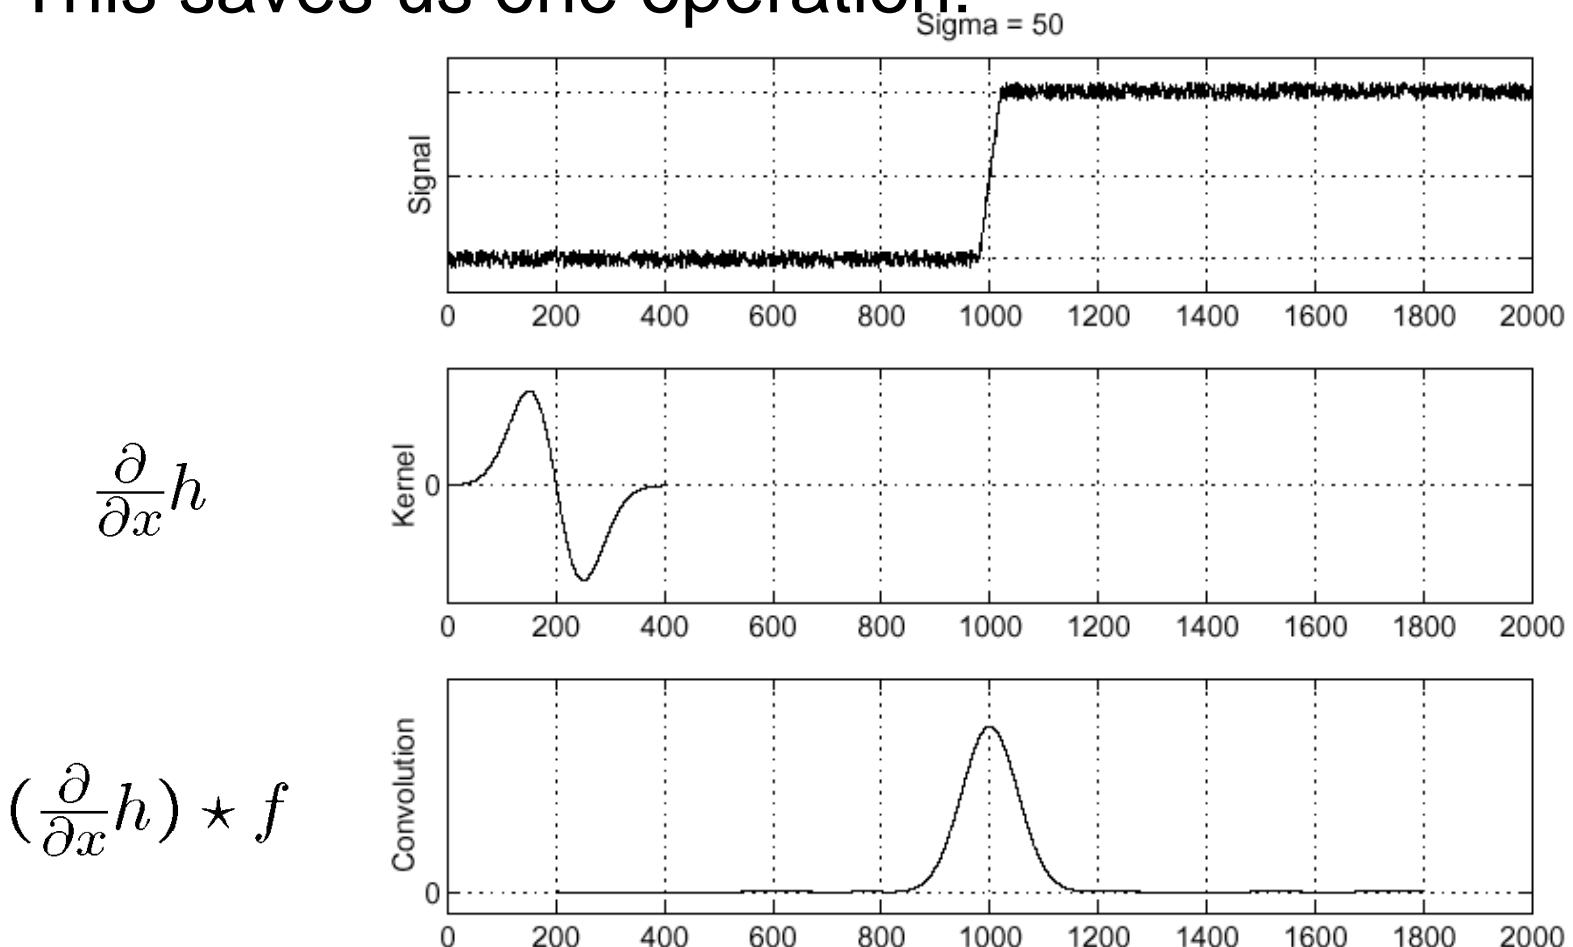
- To find edges, look for peaks in $\frac{d}{dx}(f * g)$

Source: S. Seitz

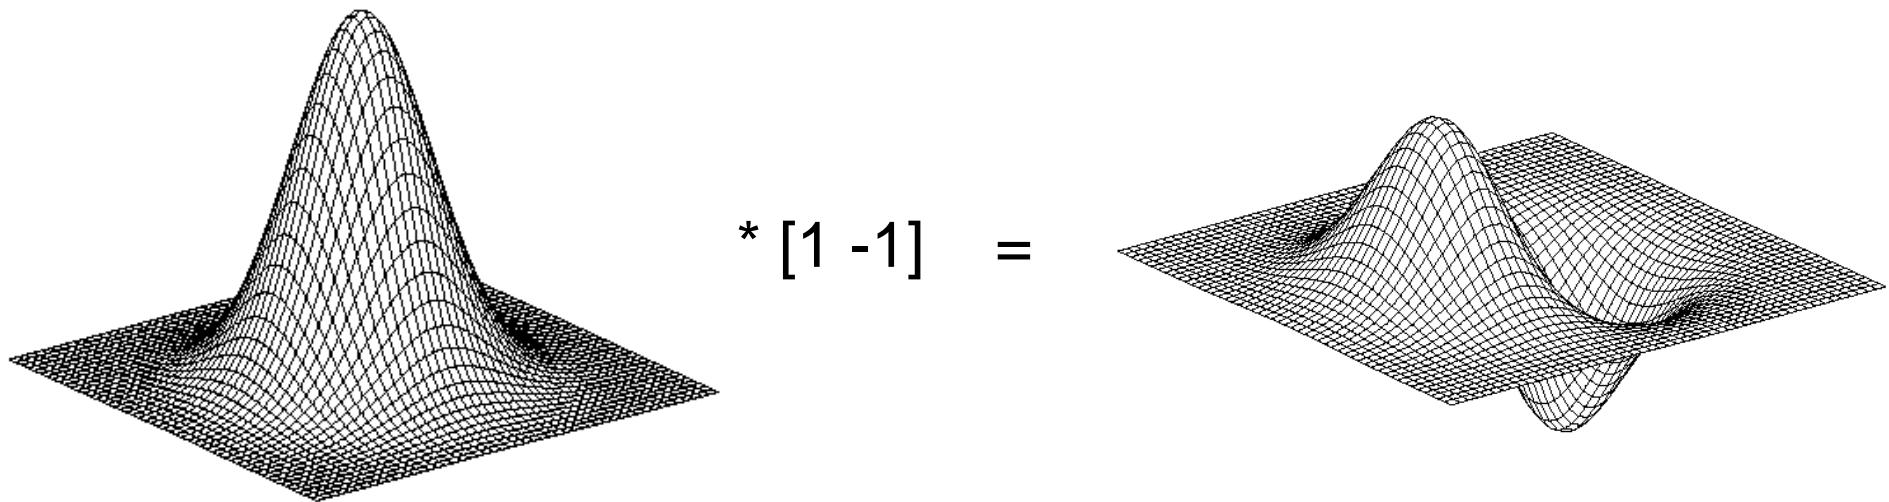
Derivative theorem of convolution

$$\frac{\partial}{\partial x}(h * f) = (\frac{\partial}{\partial x}h) * f$$

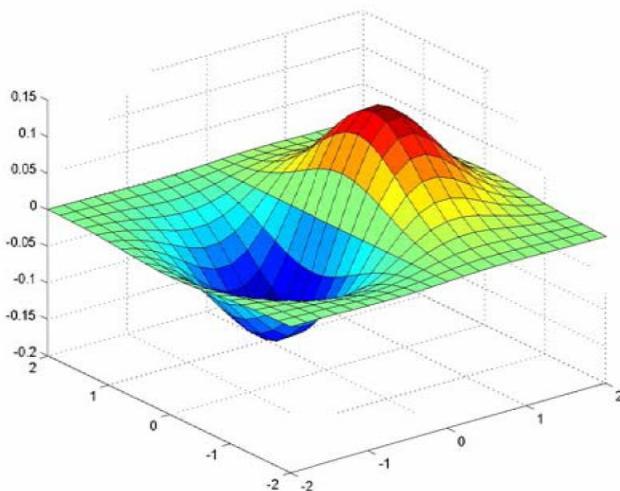
This saves us one operation:



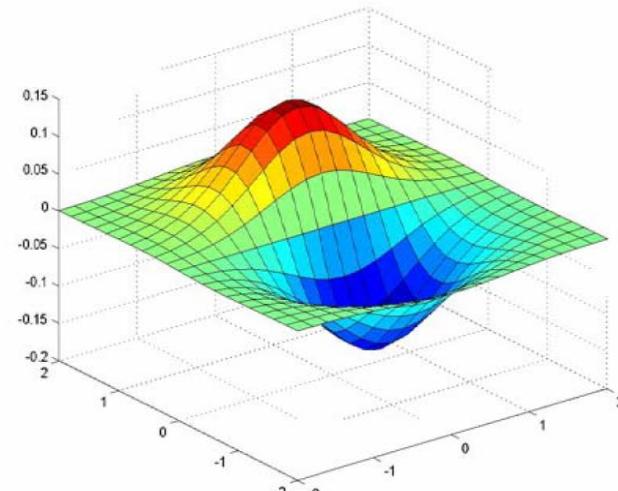
Derivative of Gaussian filter



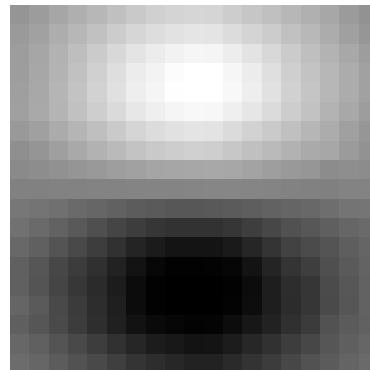
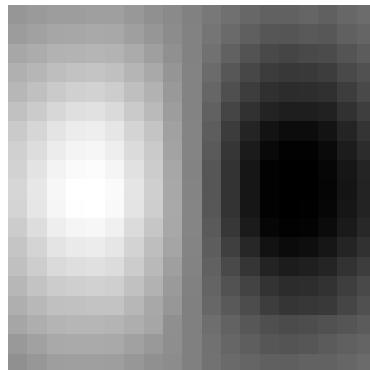
Derivative of Gaussian filter



x-direction



y-direction



Which one finds horizontal/vertical edges?

Example



input image (“Lena”)

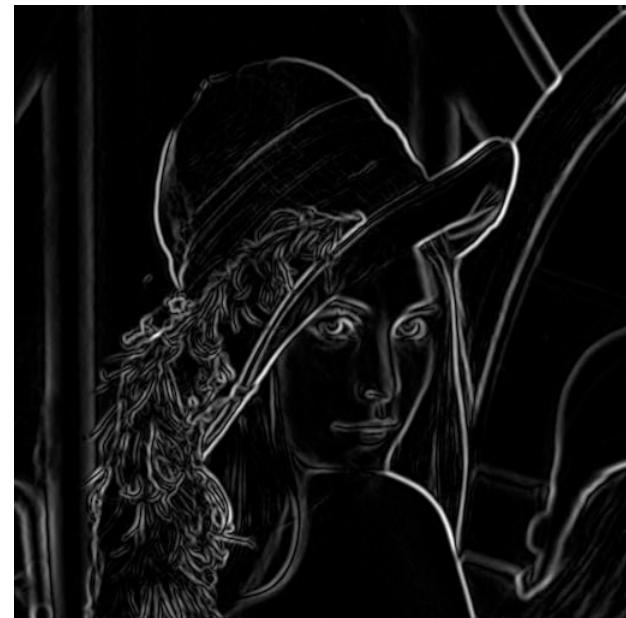
Compute Gradients (DoG)



X-Derivative of
Gaussian



Y-Derivative of
Gaussian

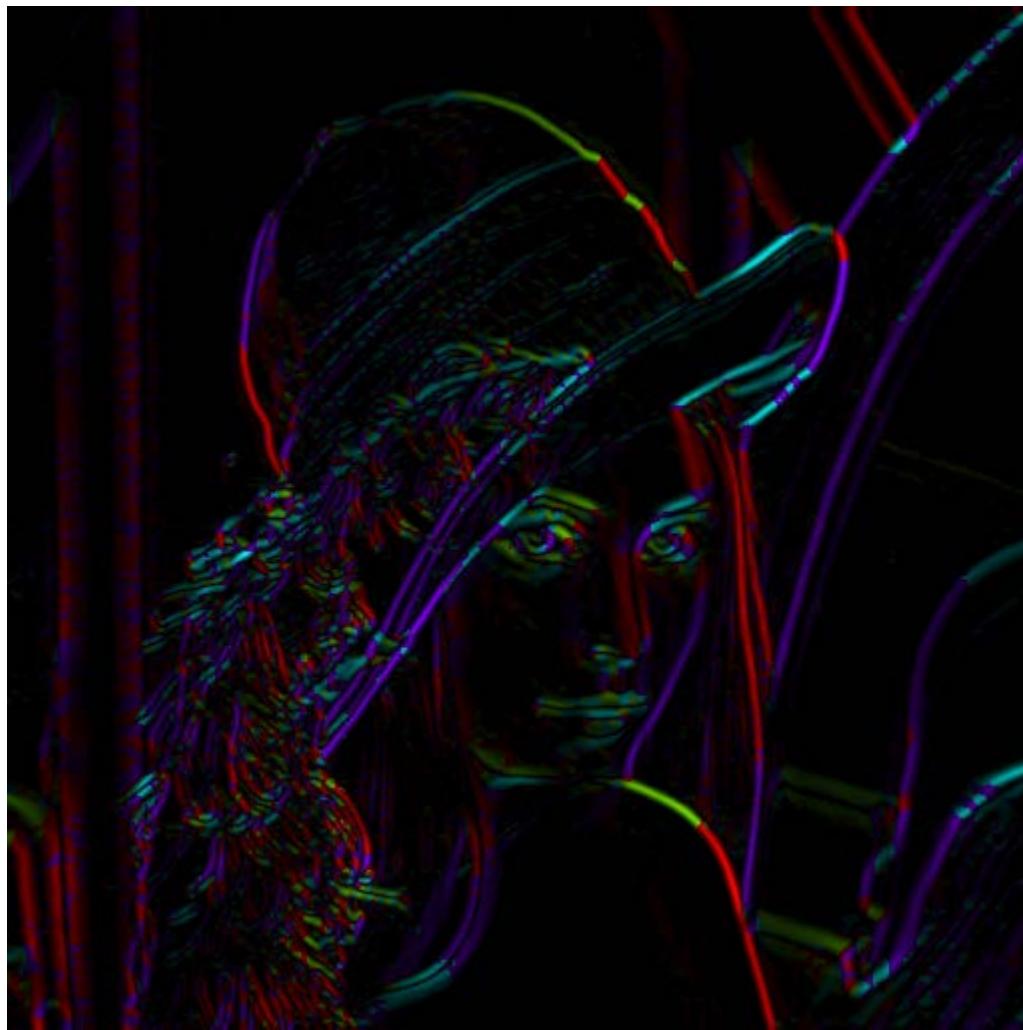


Gradient Magnitude

Get Orientation at Each Pixel

Threshold at minimum level

Get orientation



$\theta = \text{atan2}(-gy, gx)$

MATLAB demo

```
im = im2double(imread(filemane));
g = fspecial('gaussian',15,2);
imagesc(g);
surf1(g);
gim = conv2(im,g,'same');
imagesc(conv2(im,[-1 1],'same'));
imagesc(conv2(gim,[-1 1],'same')) ;
dx = conv2(g,[-1 1],'same');
Surf1(dx);
imagesc(conv2(im,dx,'same')) ;
```

Finite difference filters

Other approximations of derivative filters exist:

Prewitt: $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$; $M_y = \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$

Sobel: $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$; $M_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$

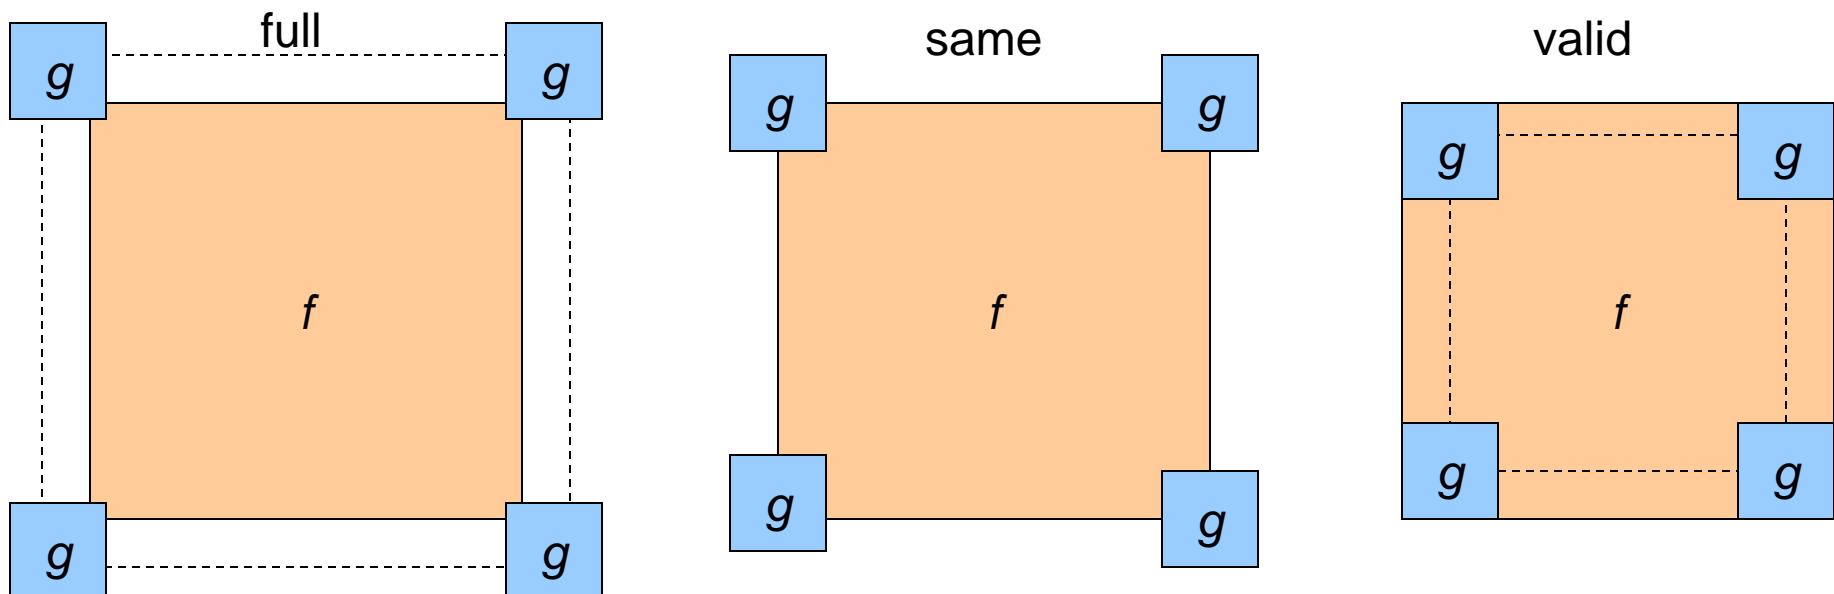
Roberts: $M_x = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$; $M_y = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$

Practical matters

What is the size of the output?

MATLAB: `filter2(g, f, shape)` or `conv2(g,f,shape)`

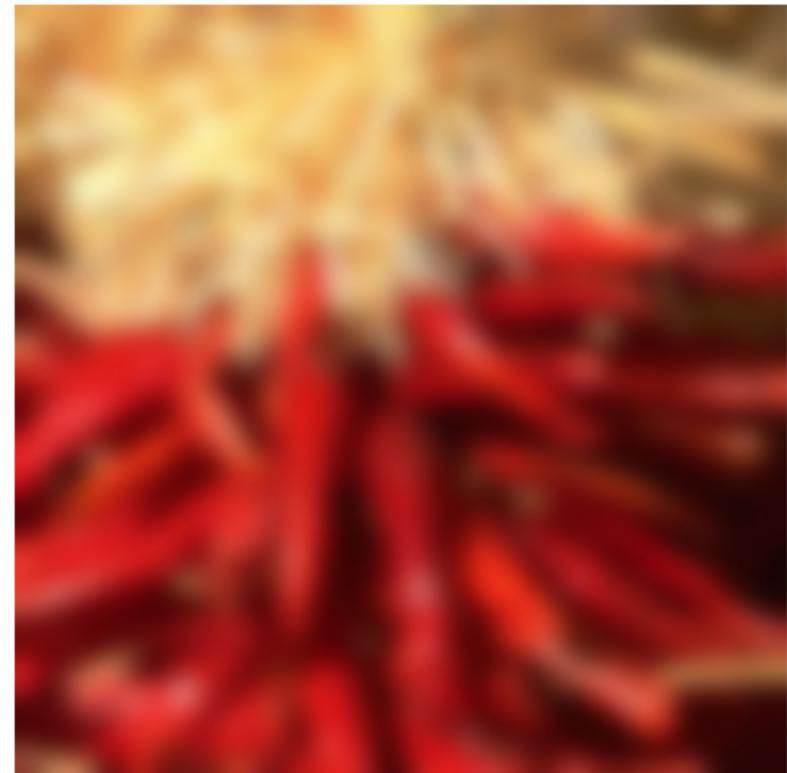
- *shape = ‘full’*: output size is sum of sizes of *f* and *g*
- *shape = ‘same’*: output size is same as *f*
- *shape = ‘valid’*: output size is difference of sizes of *f* and *g*



Practical matters

What about near the edge?

- the filter window falls off the edge of the image
- need to extrapolate
- methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge



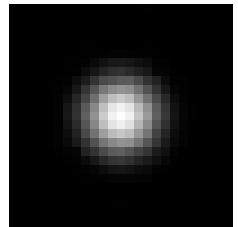
Practical matters

- methods (MATLAB):
 - clip filter (black): `imfilter(f, g, 0)`
 - wrap around: `imfilter(f, g, 'circular')`
 - copy edge: `imfilter(f, g, 'replicate')`
 - reflect across edge: `imfilter(f, g, 'symmetric')`

Review: Smoothing vs. derivative filters

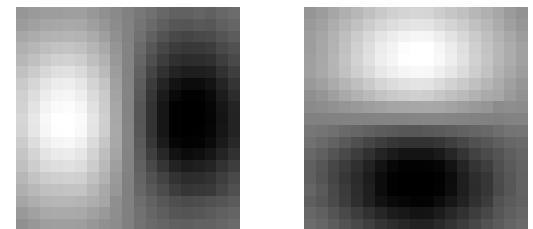
Smoothing filters

- Gaussian: remove “high-frequency” components; “low-pass” filter
- Can the values of a smoothing filter be negative?
- What should the values sum to?
 - **One**: constant regions are not affected by the filter



Derivative filters

- Derivatives of Gaussian
- Can the values of a derivative filter be negative?
- What should the values sum to?
 - **Zero**: no response in constant regions
- High absolute value at points of high contrast

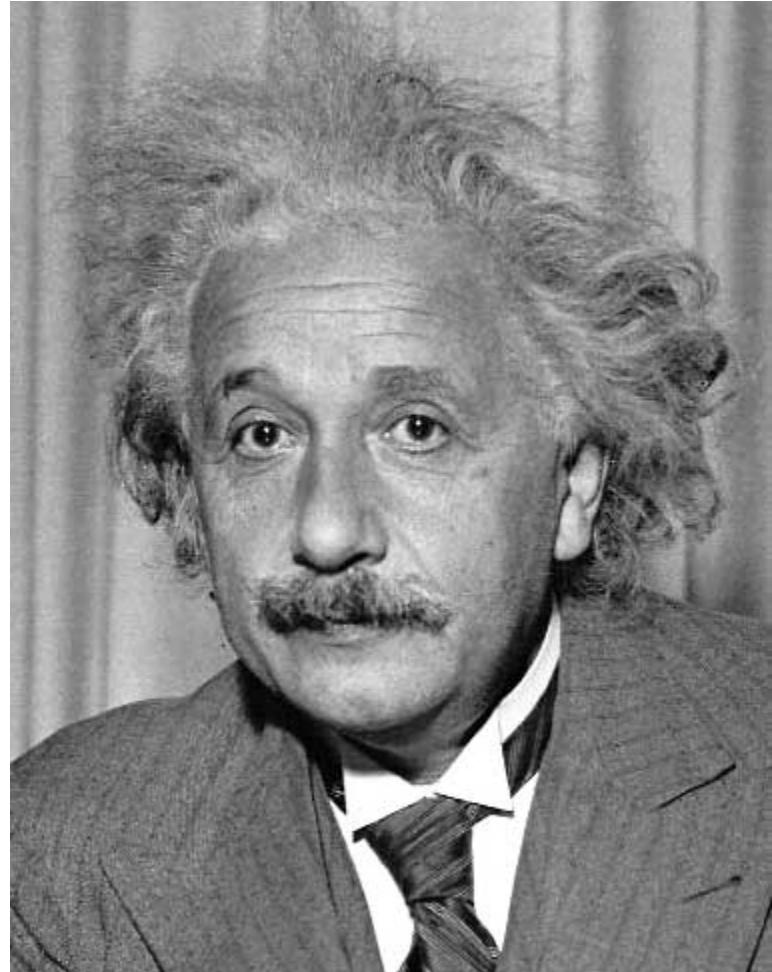


Template matching

Goal: find  in image

Main challenge: What is a good similarity or distance measure between two patches?

- Correlation
- Zero-mean correlation
- Sum Square Difference
- Normalized Cross Correlation

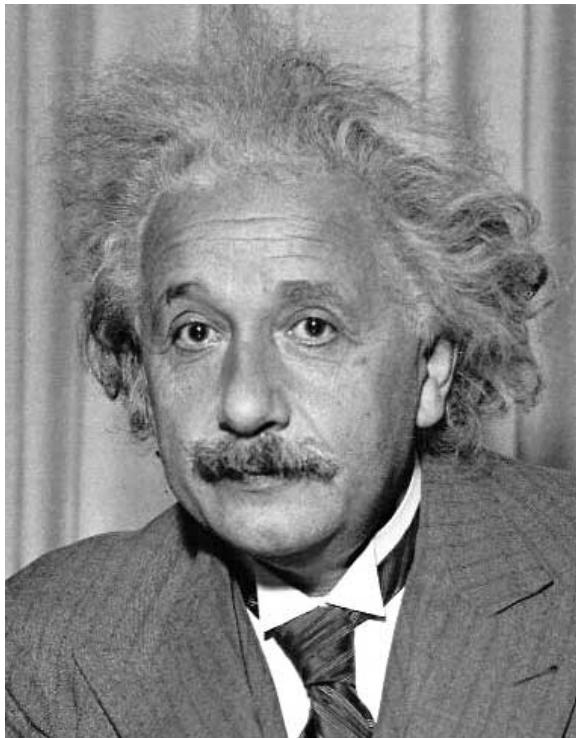


Matching with filters

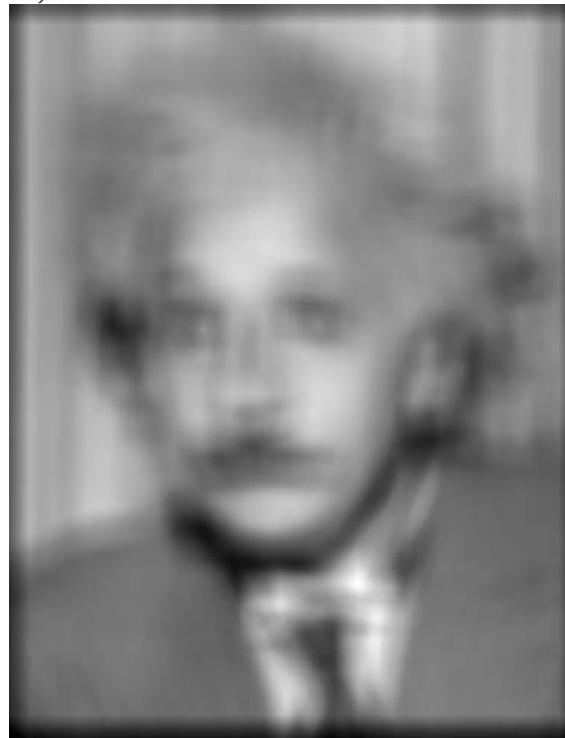
Goal: find  in image

Method 0: filter the image with eye patch

$$h[m, n] = \sum_{k, l} g[k, l] f[m + k, n + l]$$



Input



Filtered Image

f = image
g = filter

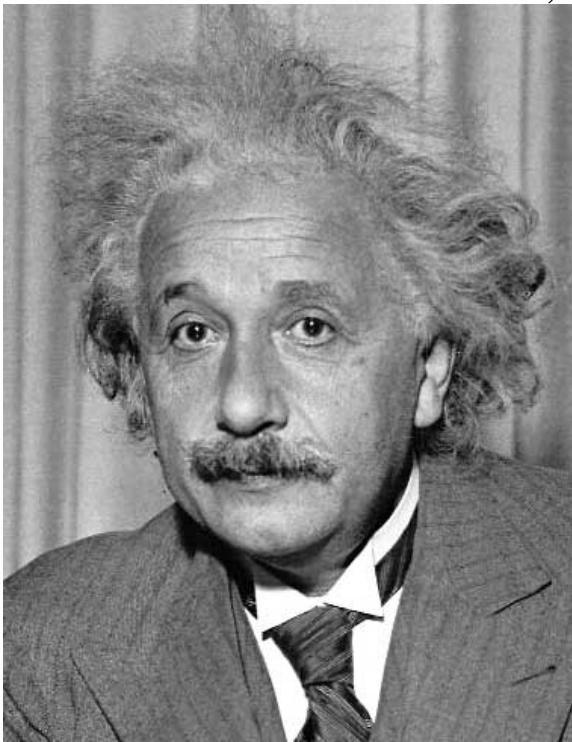
What went wrong?

Matching with filters

Goal: find  in image

Method 1: filter the image with zero-mean eye

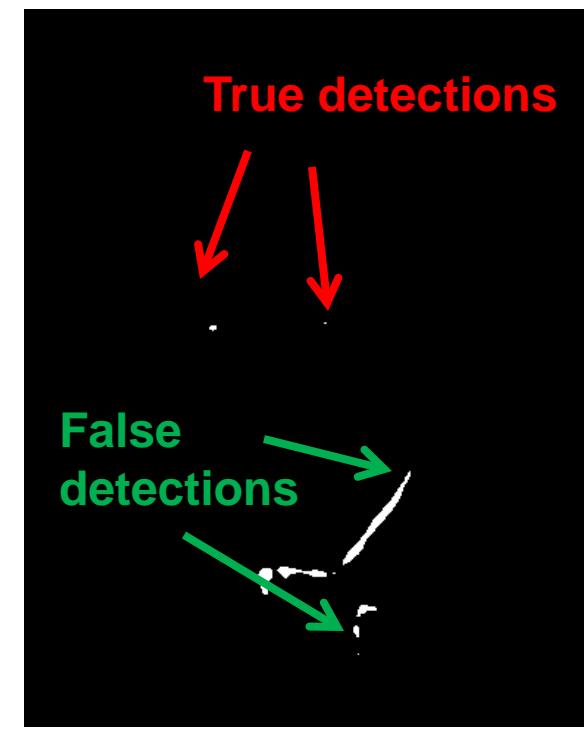
$$h[m, n] = \sum_{k, l} (f[k, l] - \bar{f}) \underbrace{(g[m + k, n + l])}_{\text{mean of } f}$$



Input



Filtered Image (scaled)



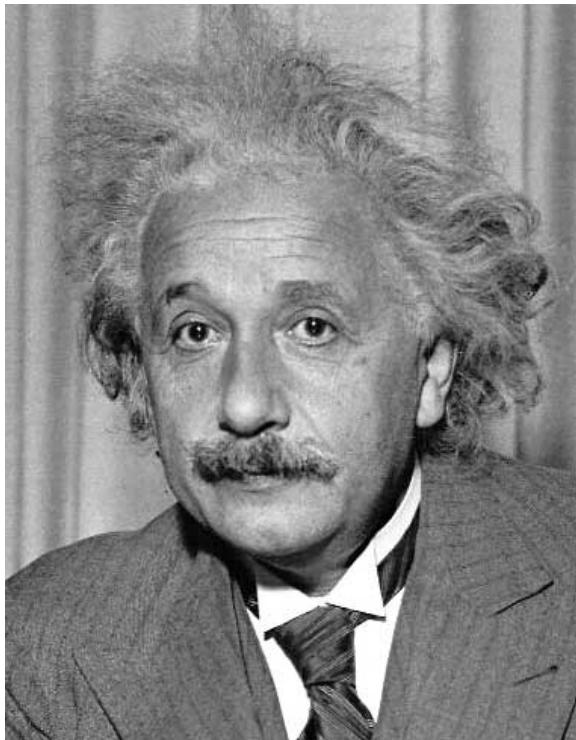
Thresholded Image

Matching with filters

Goal: find  in image

Method 2: SSD

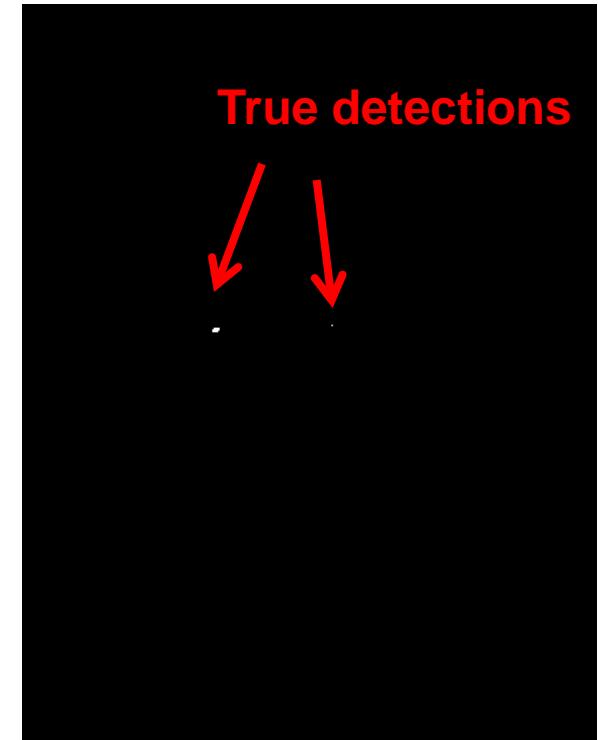
$$h[m, n] = \sum_{k, l} (g[k, l] - f[m + k, n + l])^2$$



Input



1 - $\text{sqrt}(\text{SSD})$



Thresholded Image

True detections

Matching with filters

Can SSD be implemented with linear filters?

$$h[m, n] = \sum_{k,l} (g[k, l] - f[m + k, n + l])^2$$

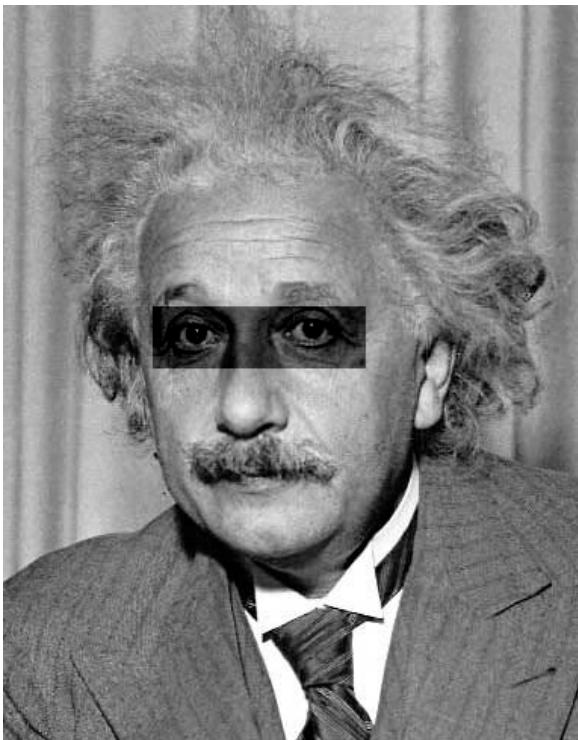
Matching with filters

Goal: find  in image

What's the potential downside of SSD?

Method 2: SSD

$$h[m,n] = \sum_{k,l} (g[k,l] - f[m+k, n+l])^2$$



Input



1- sqrt(SSD)

Side by Derek Hoiem

Matching with filters

Goal: find  in image

Method 3: Normalized cross-correlation

$$h[m, n] = \frac{\sum_{k,l} (g[k, l] - \bar{g})(f[m + k, n + l] - \bar{f}_{m,n})}{\left(\sum_{k,l} (g[k, l] - \bar{g})^2 \sum_{k,l} (f[m + k, n + l] - \bar{f}_{m,n})^2 \right)^{0.5}}$$

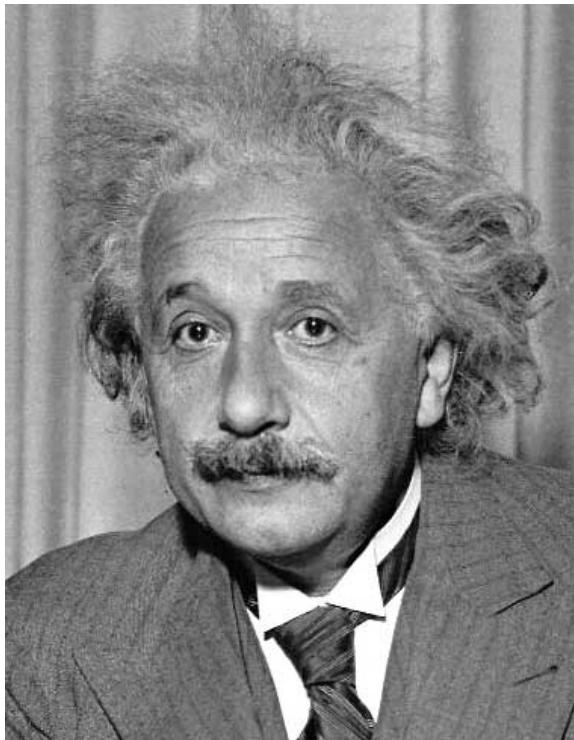
mean template
↓
 $\sum_{k,l} (g[k, l] - \bar{g})(f[m + k, n + l] - \bar{f}_{m,n})$

mean image patch
↓

Matching with filters

Goal: find  in image

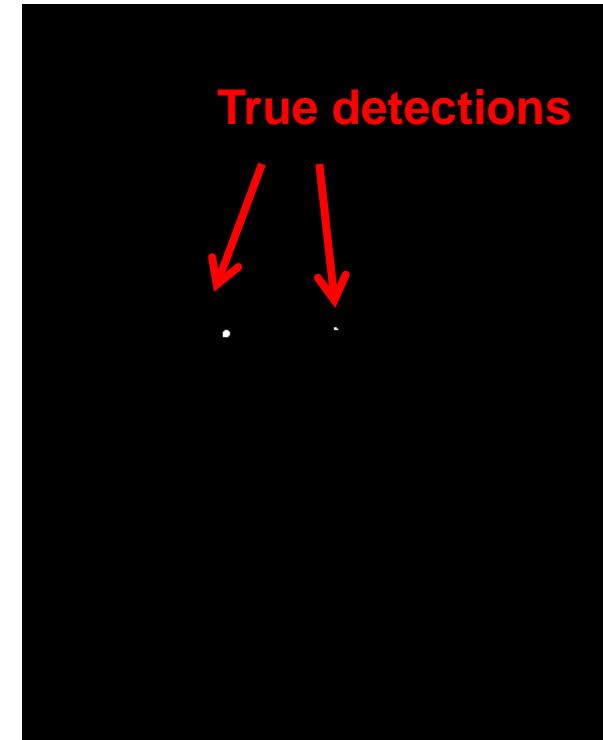
Method 3: Normalized cross-correlation



Input



Normalized X-Correlation

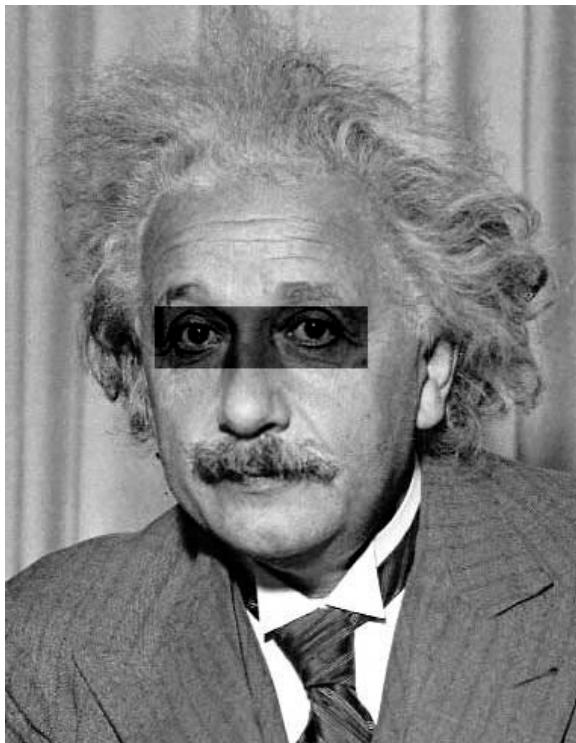


True detections
Thresholded Image

Matching with filters

Goal: find  in image

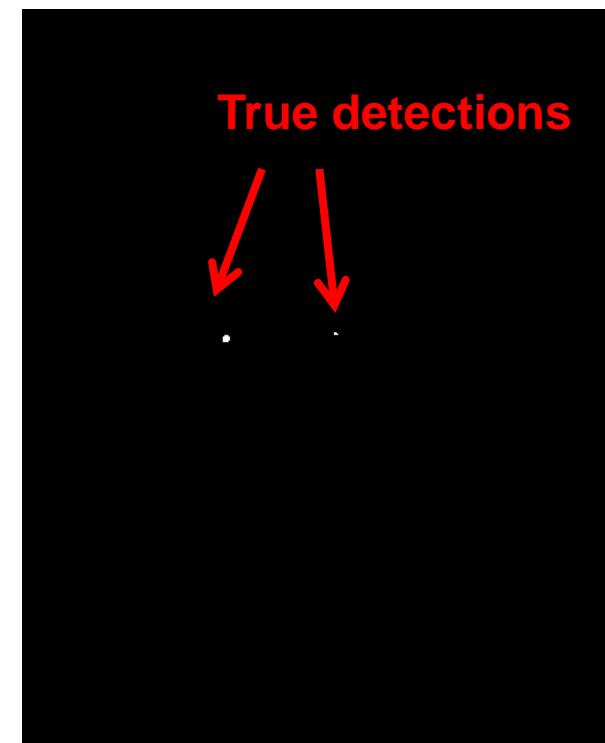
Method 3: Normalized cross-correlation



Input



Normalized X-Correlation



True detections
↓
↓
Thresholded Image

Q: What is the best method to use?

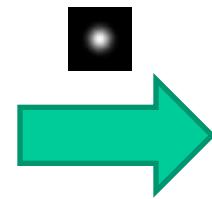
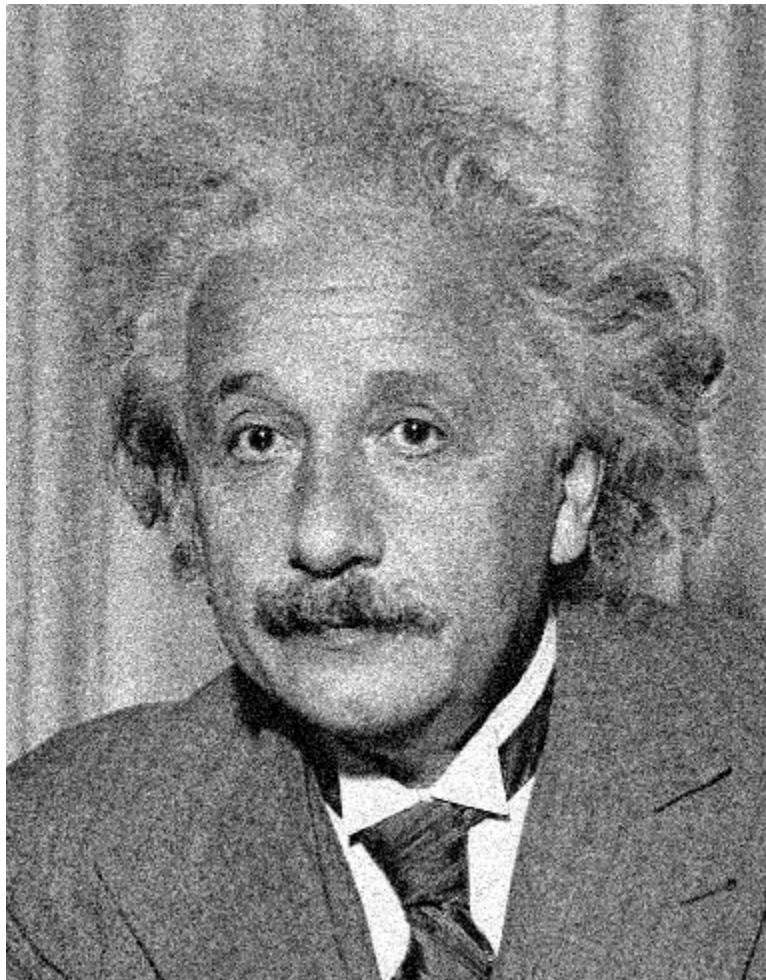
A: Depends

Zero-mean filter: fastest but not a great matcher

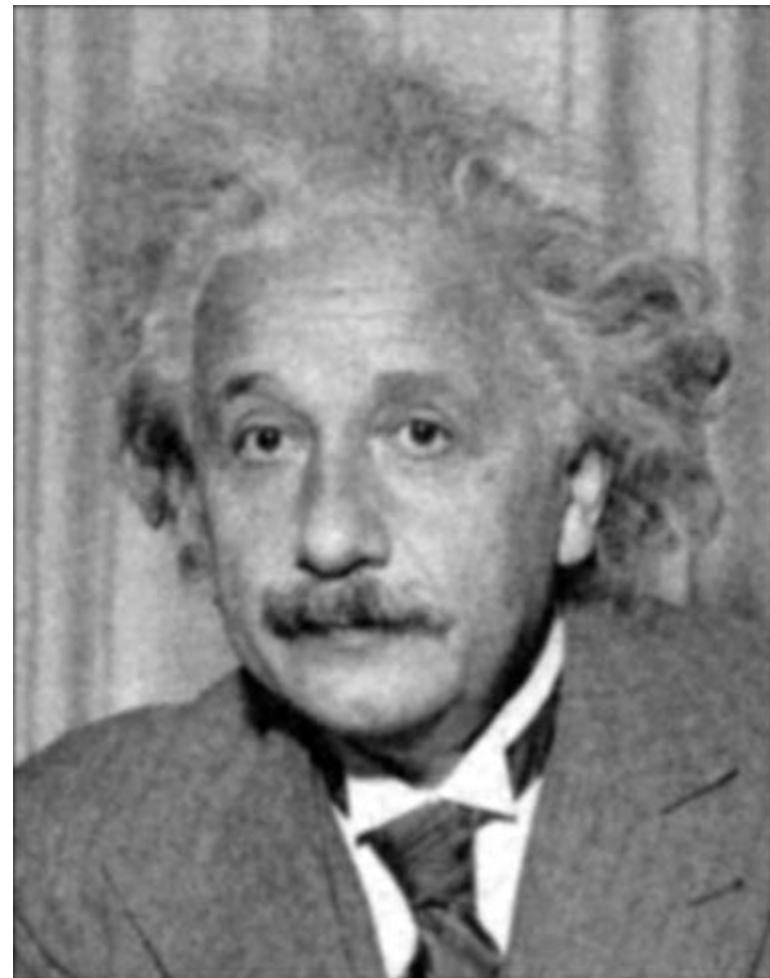
SSD: next fastest, sensitive to overall intensity

Normalized cross-correlation: slowest, invariant to local average intensity and contrast

Denoising



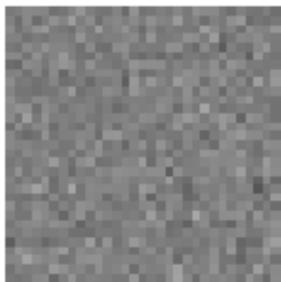
Gaussian
Filter



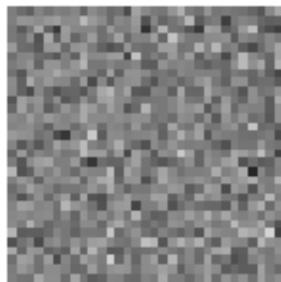
Additive Gaussian Noise

Reducing Gaussian noise

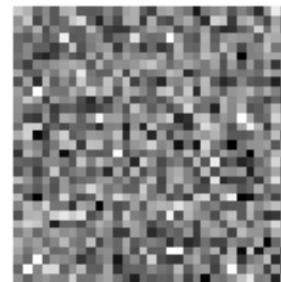
$\sigma=0.05$



$\sigma=0.1$



$\sigma=0.2$



no
smoothing



$\sigma=1$ pixel



$\sigma=2$ pixels



Smoothing with larger standard deviations suppresses noise, but also blurs the image

Reducing salt-and-pepper noise by Gaussian smoothing

3x3



5x5

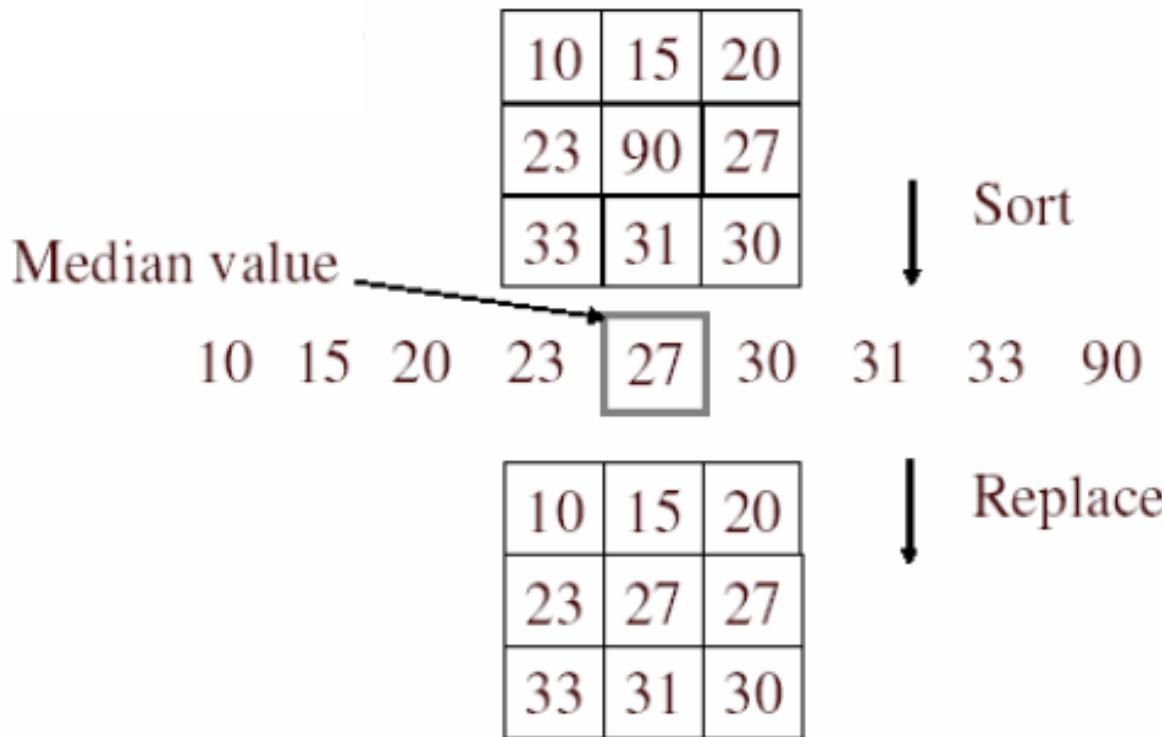


7x7



Alternative idea: Median filtering

A **median filter** operates over a window by selecting the median intensity in the window



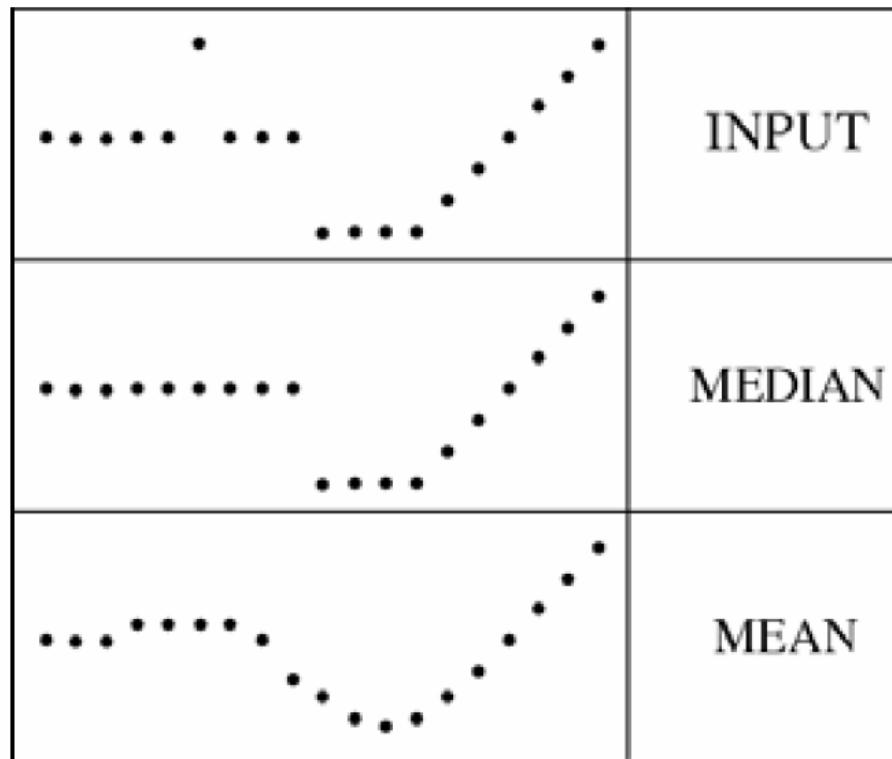
- Is median filtering linear?

Median filter

What advantage does median filtering have over Gaussian filtering?

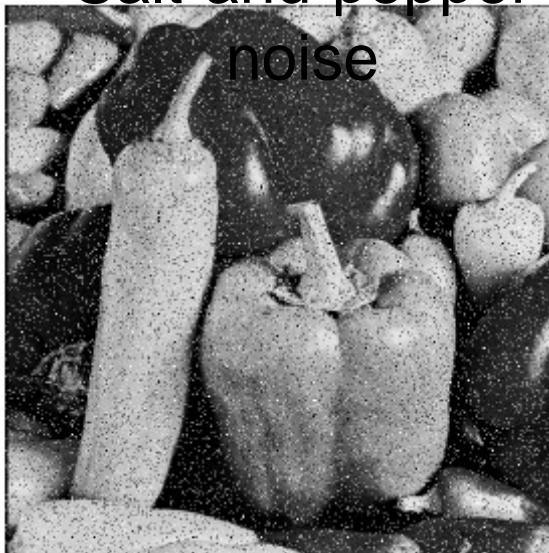
- Robustness to outliers

filters have width 5 :

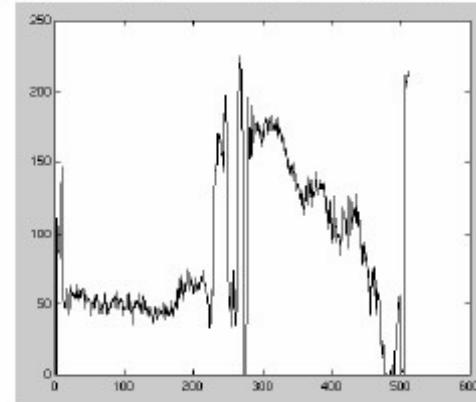
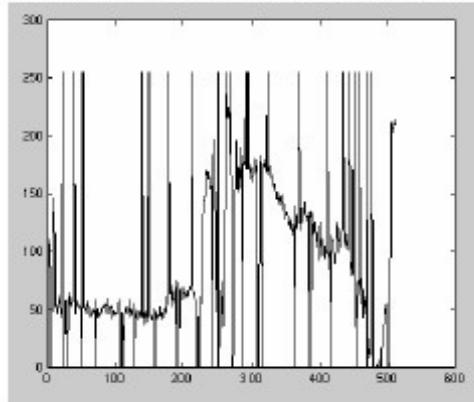


Median filter

Salt-and-pepper noise



Median filtered



MATLAB: `medfilt2(image, [h w])`

Source: M. Hebert

Median vs. Gaussian filtering

3x3



5x5



7x7



Gaussian

Median





EXTRA SLIDES

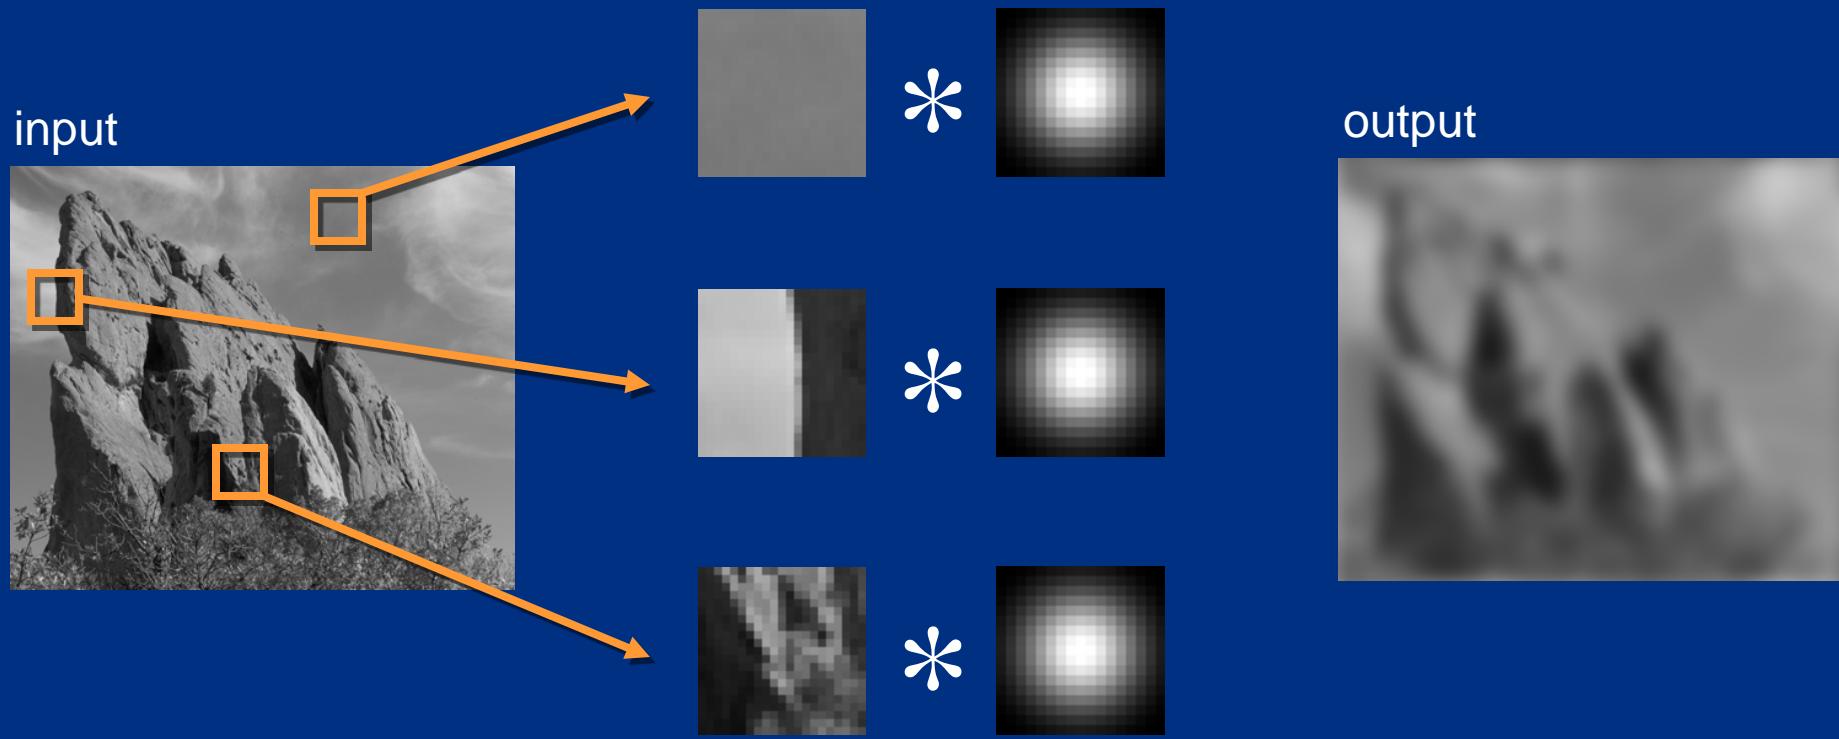
A Gentle Introduction to Bilateral Filtering and its Applications



“Fixing the Gaussian Blur”: the Bilateral Filter

Sylvain Paris – MIT CSAIL

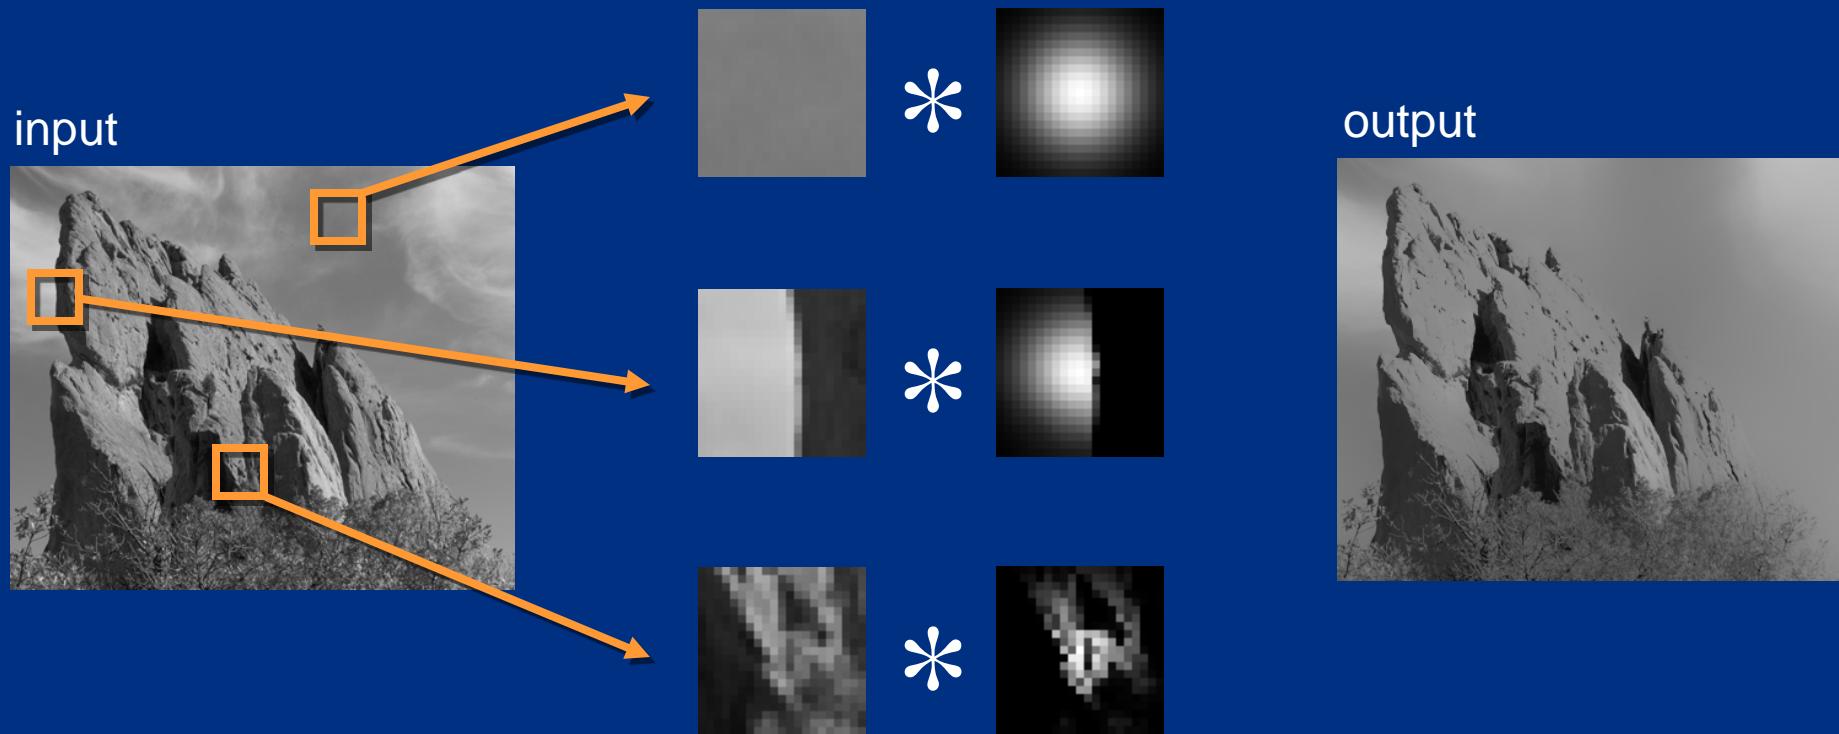
Blur Comes from Averaging across Edges



Same Gaussian kernel everywhere.

Bilateral Filter [Aurich 95, Smith 97, Tomasi 98]

No Averaging across Edges



The kernel shape depends on the image content.

Bilateral Filter Definition: an Additional Edge Term

Same idea: **weighted average of pixels.**

$$BF[I]_p = \frac{1}{W_p} \sum_{q \in S} G_{\sigma_s}(\| p - q \|) G_{\sigma_r}(|I_p - I_q|) I_q$$

new
not new
new

↓
↓
↓

normalization factor **space** weight **range** weight

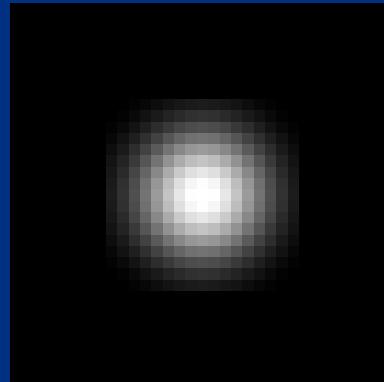
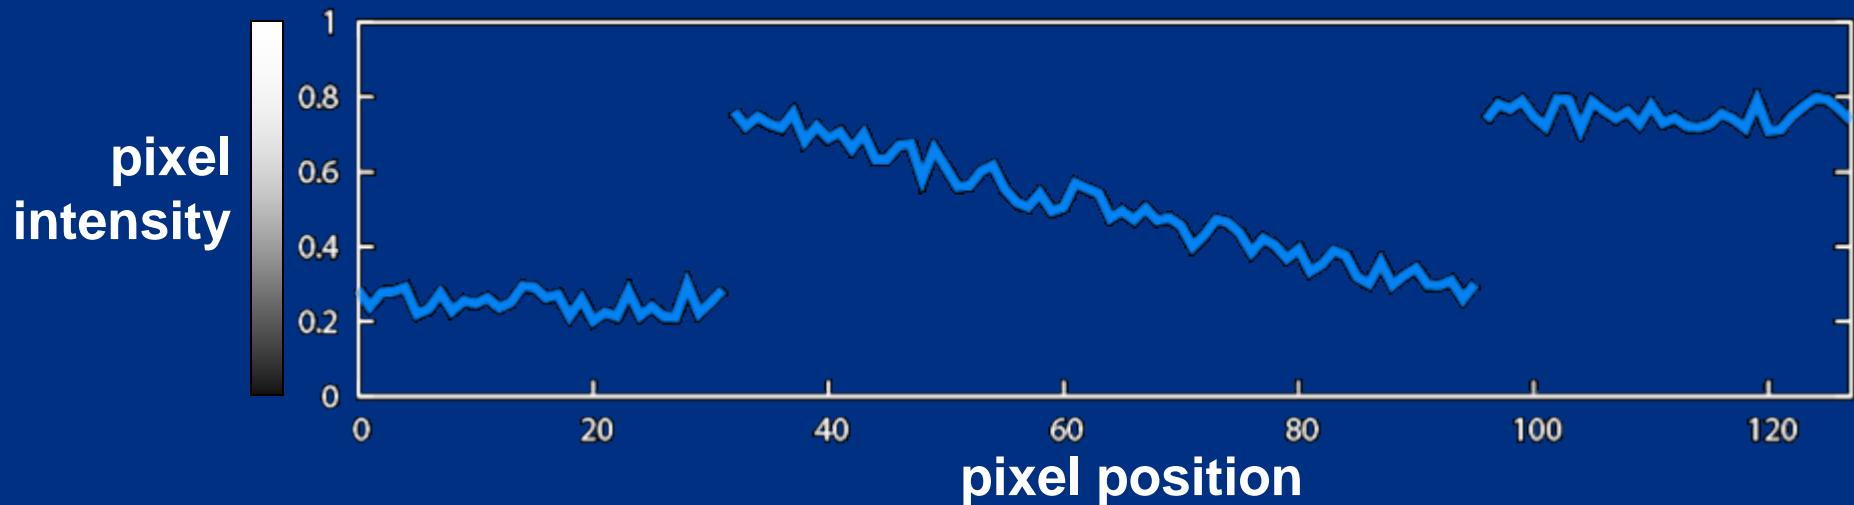


Illustration a 1D Image

- 1D image = line of pixels

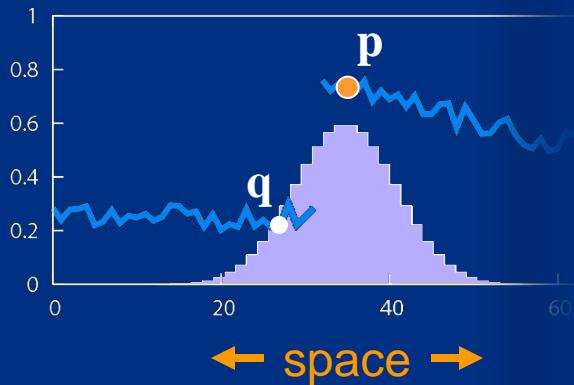


- Better visualized as a plot



Gaussian Blur and Bilateral Filter

Gaussian blur

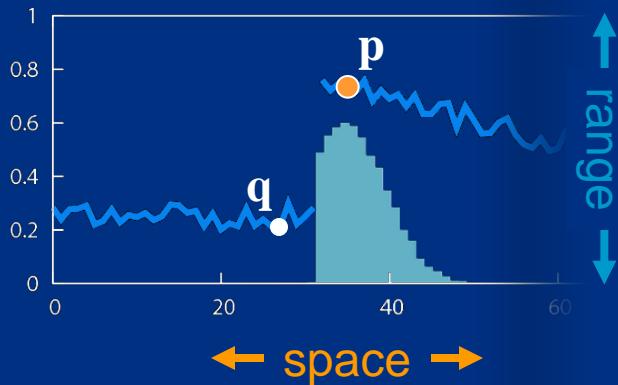


$$GB[I]_p = \sum_{q \in S} G_\sigma(\|p - q\|) I_q$$

space

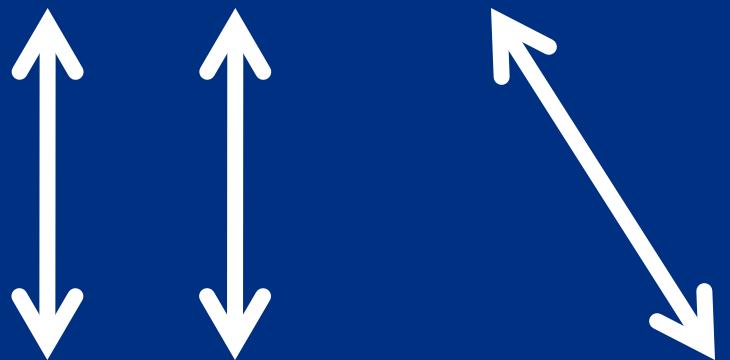
Bilateral filter

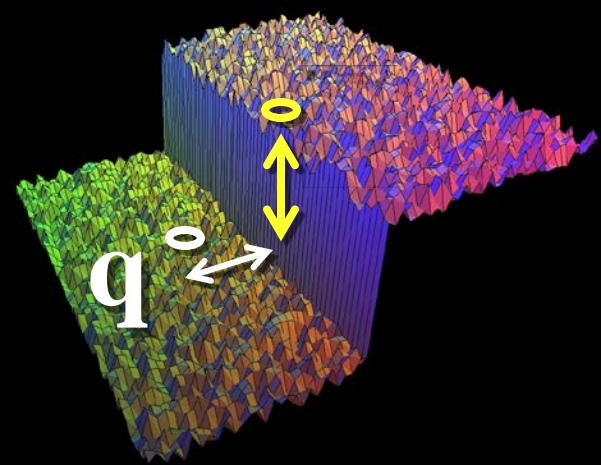
[Aurich 95, Smith 97, Tomasi 98]



$$BF[I]_p = \frac{1}{W_p} \sum_{q \in S} G_{\sigma_s}(\|p - q\|) G_{\sigma_r}(|I_p - I_q|) I_q$$

space range
normalization





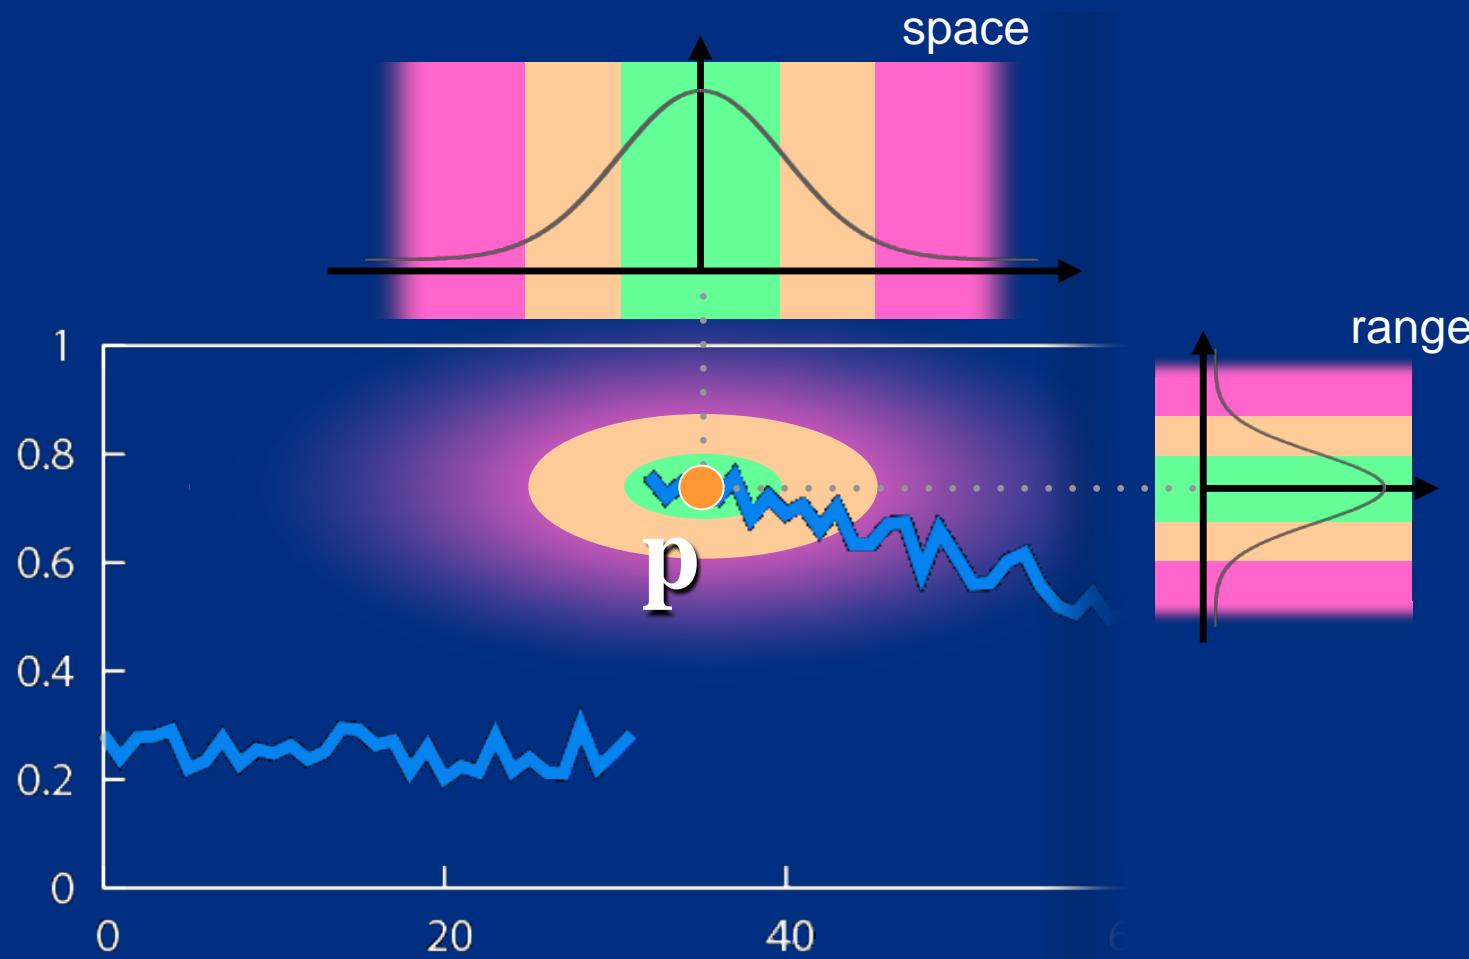
Space and Range Parameters

$$BF[I]_p = \frac{1}{W_p} \sum_{q \in S} G_{\sigma_s}(\|p - q\|) G_{\sigma_r}(|I_p - I_q|) I_q$$


- space σ_s : spatial extent of the kernel, size of the considered neighborhood.
- range σ_r : “minimum” amplitude of an edge

Influence of Pixels

Only pixels close in space and in range are considered.



Exploring the Parameter Space



input

$\sigma_s = 2$



$\sigma_r = 0.1$

$\sigma_r = 0.25$

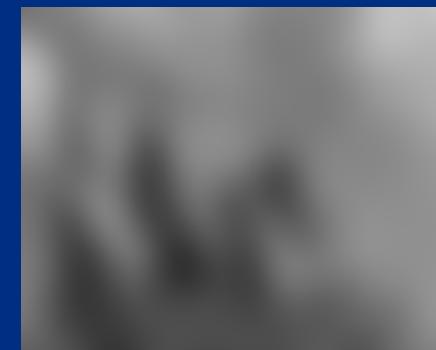
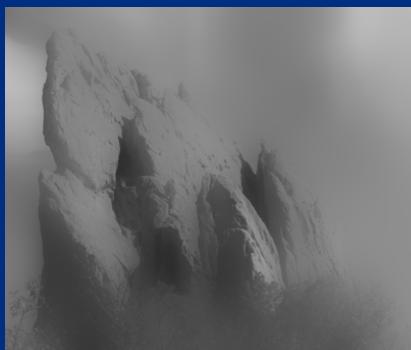
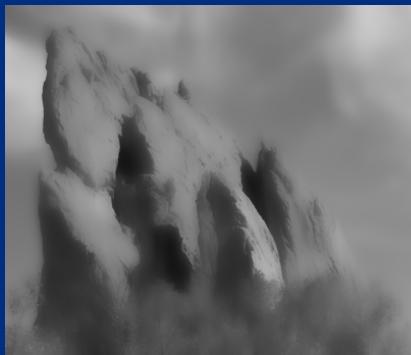
$\sigma_r = \infty$
(Gaussian blur)



$\sigma_s = 6$



$\sigma_s = 18$



Varying the Range Parameter



input

$\sigma_s = 2$



$\sigma_r = 0.1$

$\sigma_r = 0.25$

$\sigma_r = \infty$
(Gaussian blur)



$\sigma_s = 6$



$\sigma_s = 18$



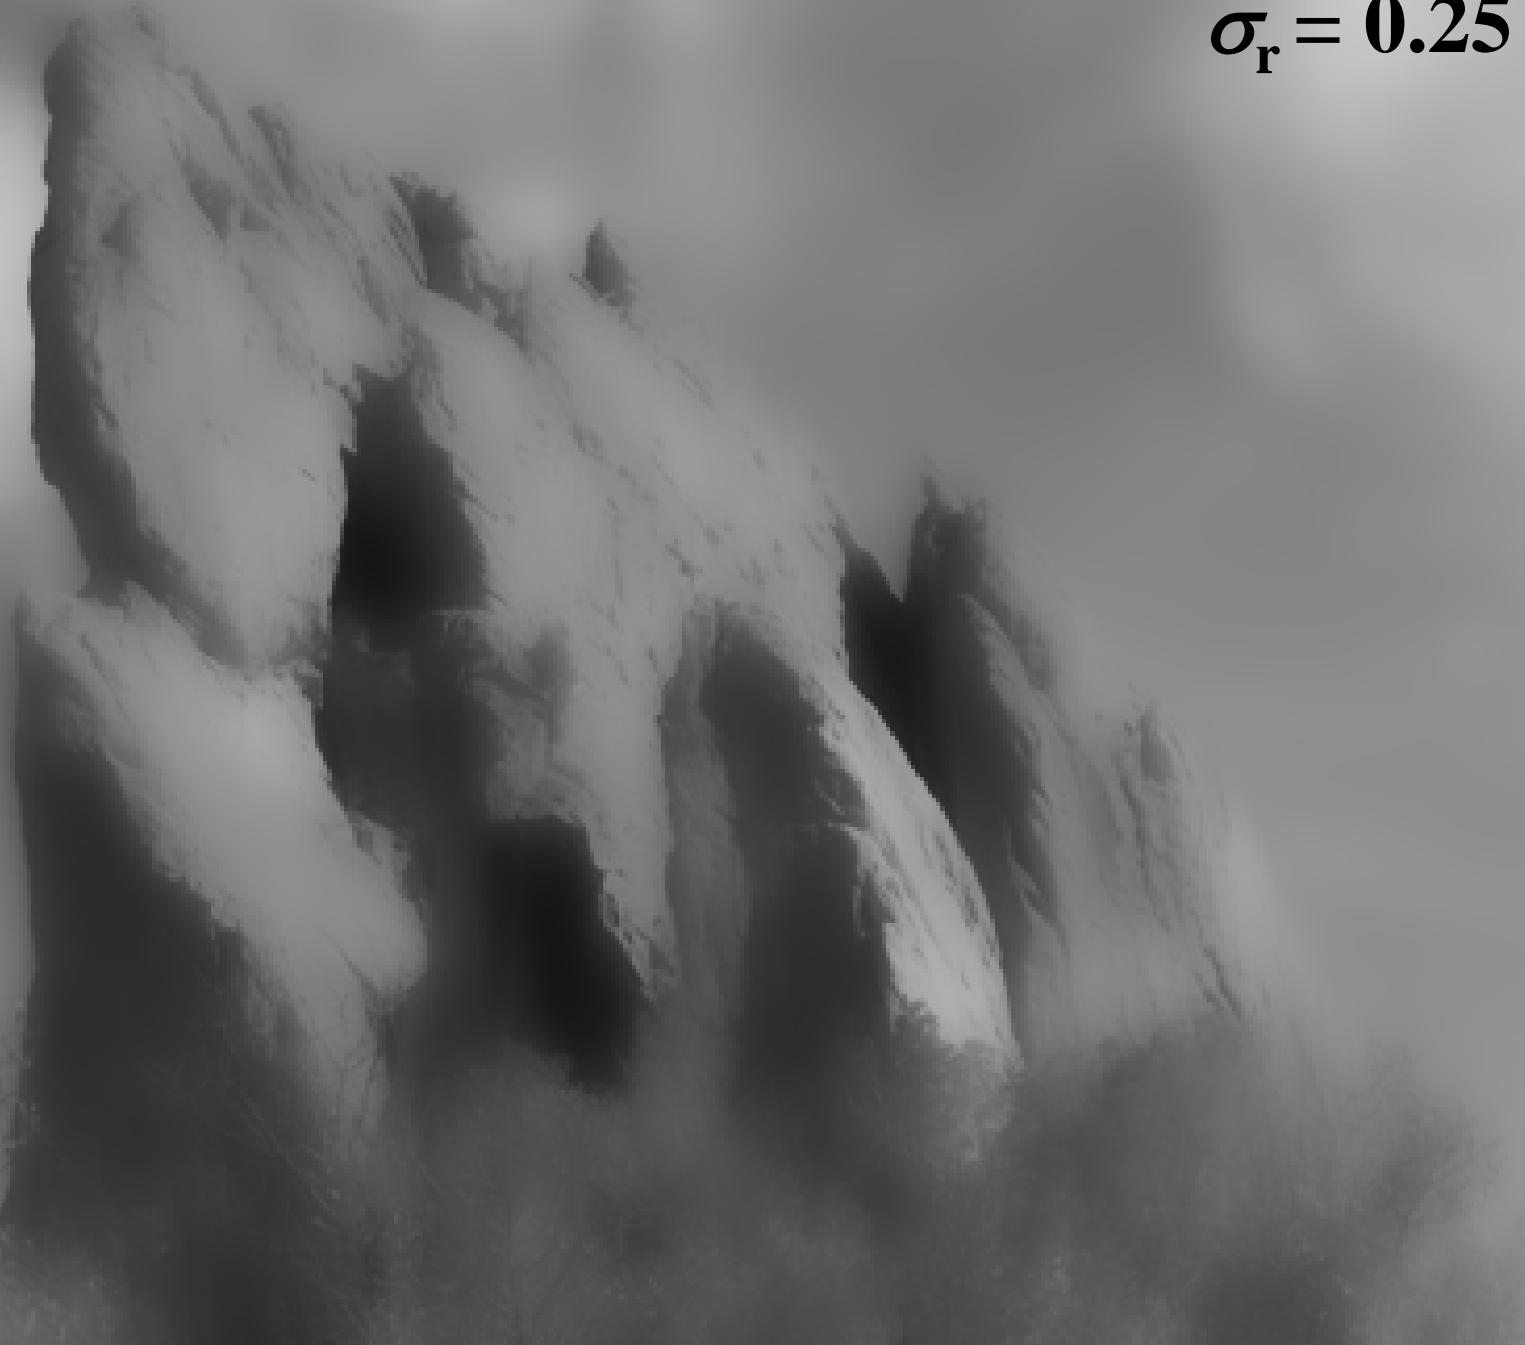
input



$\sigma_r = 0.1$



$\sigma_r = 0.25$



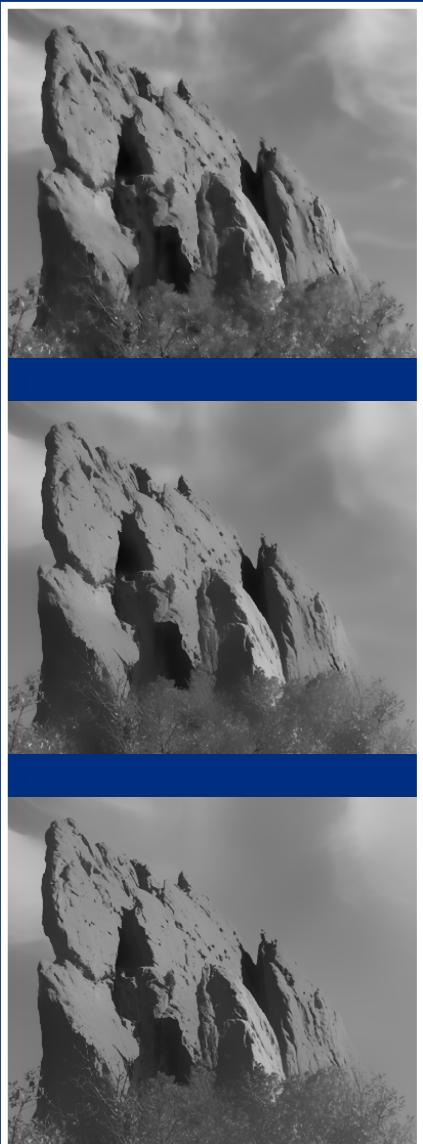
$\sigma_r = \infty$
(Gaussian blur)

Varying the Space Parameter



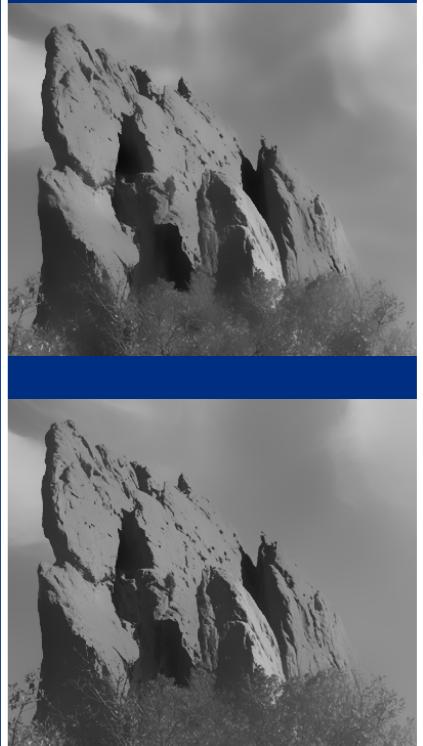
input

$\sigma_s = 2$



$\sigma_r = 0.1$

$\sigma_s = 6$



$\sigma_s = 18$



$\sigma_r = 0.25$

$\sigma_r = \infty$
(Gaussian blur)



input



$$\sigma_s = 2$$



$\sigma_s = 6$ 

$\sigma_s = 18$ 